Enhancing xv6 with a File Descriptor Information System Call

CSE323 Project Report - Group 8
Feature 2: File Descriptor Information System Call (fdinfo)

1. Project Goals

This report describes the implementation and validation of Feature 2: a File Descriptor Information system call (fdinfo) in xv6. The goal is to provide user-space processes with direct runtime visibility into their open file descriptors, active byte offsets, and underlying filesystem or pipe metadata.

In Unix-like systems, a file descriptor is an integer handle referencing an entry in the system-wide open file table. While xv6 includes the fstat() system call, fstat() only returns static disk inode metadata (such as file type, inode number, and size). It cannot expose descriptor-level runtime properties, such as the current read/write byte offset (f->off), access permissions (f->readable, f->writable), or endpoint directionality for pipes. The added fdinfo() system call connects user programs with per-process descriptor state while enforcing strict kernel memory safety.

Field	Type	Source in Kernel	Role
type	int	f->type	Descriptor type: FD_TYPE_NONE (0), FD_TYPE_PIPE (1), FD_TYPE_INODE (2).
dev	int	f->ip->dev	Disk device number containing the underlying file inode.
inum	uint	f->ip->inum	Unique inode number on the storage device.
nlink	short	f->ip->nlink	Number of active directory hard links to this file.
size	uint	f->ip->size	Total file length in bytes (from disk inode).
off	uint	f->off	Current read/write byte offset in the open file stream.
readable / writable	char	f->readable / writable	Access mode flags (1/0) granted during open() or pipe().

Concurrency and memory safety rules are enforced: the kernel validates user-space buffer pointers and safely acquires the inode sleep-lock (ilock) during metadata extraction to prevent data races.

2. Modifications

The implementation required changes to shared headers, system call dispatch tables, the user interface assembly layer, filesystem system call handlers, and user-level test suite support.

2.1 Data Structure Definitions

The shared header stat.h was updated to define struct fdinfo and descriptor type constants accessible to both kernel files and user programs.

```
// stat.h
#define FD_TYPE_NONE  0
#define FD_TYPE_PIPE  1
#define FD_TYPE_INODE 2

struct fdinfo {
  int type;          // FD_TYPE_NONE, FD_TYPE_PIPE, FD_TYPE_INODE
  int dev;           // Device number (if inode)
  uint inum;         // Inode number (if inode)
  short nlink;       // Number of links (if inode)
  uint size;         // Size of file in bytes (if inode)
  uint off;          // Current file offset
  char readable;     // 1 if readable, 0 otherwise
  char writable;     // 1 if writable, 0 otherwise
};
```

2.2 System Call Wiring

The new system call was registered with number 23 in syscall.h, mapped to sys_fdinfo in syscall.c, declared in user.h, and assigned a trap entry in usys.S.

```
// syscall.h
#define SYS_fdinfo 23

// syscall.c
extern int sys_fdinfo(void);
[SYS_fdinfo]  sys_fdinfo,

// user.h
struct fdinfo;
int fdinfo(int fd, struct fdinfo*);

// usys.S
SYSCALL(fdinfo)
```

2.3 Kernel System Call Handler

The system call handler sys_fdinfo() is implemented in sysfile.c. It uses argfd() to validate the descriptor and retrieve the struct file pointer, and argptr() to verify that the destination buffer is in valid user memory.

```
// sysfile.c, inside sys_fdinfo()
int
sys_fdinfo(void)
{
  int fd;
  struct file *f;
  struct fdinfo *info;

  if(argfd(0, &fd, &f) < 0 || argptr(1, (char**)&info, sizeof(*info)) < 0)
    return -1;

  memset(info, 0, sizeof(*info));
  info->readable = f->readable;
  info->writable = f->writable;

  if(f->type == FD_INODE){
    info->type = FD_TYPE_INODE;
    ilock(f->ip);
    info->dev = f->ip->dev;
    info->inum = f->ip->inum;
    info->nlink = f->ip->nlink;
    info->size = f->ip->size;
    info->off = f->off;
    iunlock(f->ip);
  } else if(f->type == FD_PIPE){
    info->type = FD_TYPE_PIPE;
  } else {
    info->type = FD_TYPE_NONE;
  }

  return 0;
}
```

2.4 Inode Locking & Concurrency Synchronization

When inspecting an inode-backed file, the kernel acquires the inode sleep-lock via ilock(f->ip) before reading metadata (size, inum, nlink, dev) and releases it via iunlock(f->ip). This ensures race-free metadata inspection even when concurrent processes write to the same file.

2.5 In-Depth Analysis: Behavior Across Descriptor Types

The fdinfo() system call demonstrates distinct operational behaviors depending on descriptor classification:

• Regular Inode Files: Maintains stateful file offset tracking (f->off). When a process writes data, the offset advances linearly with byte count. When reopened with O_RDONLY, the offset resets to 0 while total file size remains persistent on disk.

• Pipe Endpoints: Pipes do not maintain disk offsets. Instead, fdinfo() identifies FD_TYPE_PIPE and validates unidirectional channel constraints (e.g., read-end has readable=1, writable=0; write-end has readable=0, writable=1).

• Standard Device Descriptors: File descriptors 0, 1, and 2 are categorized as device inodes (major device 1, inode 1) with size 0, confirming active binding to the console driver.

3. Evaluation

The system call was evaluated with a dedicated user-level test program named fdinfotest.c. The program tests four scenarios: standard console descriptors, dynamic file lifecycle and offset tracking across writes and reads, pipe read/write endpoint differentiation, and error rejection on invalid descriptors.

```c
// fdinfotest.c
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"

int
main(void)
{
  struct fdinfo info;
  int fd, pfd[2];

  // Test 1: Standard FDs (stdin/stdout)
  fdinfo(0, &info);
  fdinfo(1, &info);

  // Test 2: File Creation, Write, and Offset Tracking
  fd = open("fdinfo_demo.txt", O_CREATE | O_RDWR);
  fdinfo(fd, &info); // off = 0, size = 0
  write(fd, "Operating Systems CSE323!", 25);
  fdinfo(fd, &info); // off = 25, size = 25
  close(fd);

  // Reopen read-only and read 10 bytes
  fd = open("fdinfo_demo.txt", O_RDONLY);
  char buf[16];
  read(fd, buf, 10);
  fdinfo(fd, &info); // off = 10, size = 25, readable = 1, writable = 0
  close(fd);

  // Test 3: Pipe Descriptors
  pipe(pfd);
  fdinfo(pfd[0], &info); // type = PIPE, readable = 1, writable = 0
  fdinfo(pfd[1], &info); // type = PIPE, readable = 0, writable = 1
  close(pfd[0]); close(pfd[1]);

  // Test 4: Error Handling
  // fdinfo(-1, &info) -> returns -1
  // fdinfo(99, &info) -> returns -1
  // fdinfo(1, NULL)   -> returns -1
  exit();
}
```

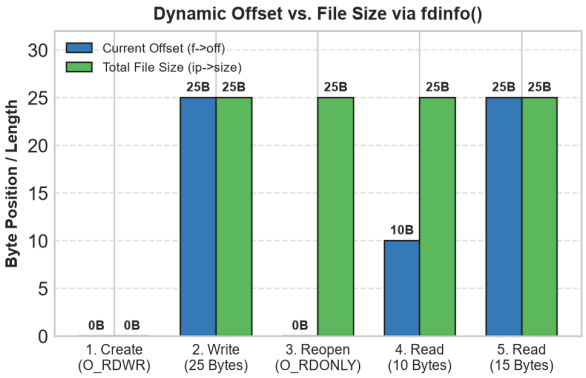

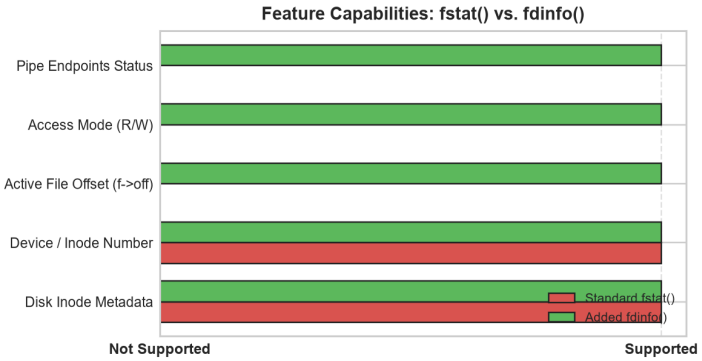

Observed Output	Meaning
FD 0 / FD 1 -> Type: INODE, dev: 1, inum: 1	Standard console descriptors are correctly recognized as device inodes.
Initial file -> Type: INODE, Size: 0, Offset: 0	Newly created file starts with 0 length and 0 byte offset.
After write 25B -> Size: 25, Offset: 25	Offset and file size dynamically increment after write operations.
Reopen O_RDONLY + read 10B -> Offset: 10, Size: 25	Offset advances to 10 bytes; read-only mode verified (r=1, w=0).
pfd[0] -> PIPE (r=1, w=0); pfd[1] -> PIPE (r=0, w=1)	Pipe endpoints are properly identified with unidirectional permissions.
fdinfo(-1), fdinfo(99), fdinfo(1, NULL) -> -1	Invalid descriptors and invalid memory pointers are safely rejected.

These outcomes demonstrate the required behavior: fdinfo() accurately exposes per-process descriptor state, dynamically reflects offset progression across file operations, differentiates pipes and devices, and prevents invalid memory access faults.

4. Conclusions

Feature 2 was completed by implementing the fdinfo() system call in xv6. The feature bridges user-space processes with kernel-level open file objects, tracking runtime byte offsets, access permissions, and underlying inode/pipe metadata. Strict synchronization via ilock/iunlock and pointer checks ensure stability and concurrency safety.

For final submission, this report should be submitted with diff_report.txt. The diff report must compare the original and modified xv6 folders and include only .c and .h file differences, as required by the project guideline.

Appendix: Diff Report Command

The required diff report can be generated from a temporary comparison repository after placing xv6-original and xv6-modified in the same parent directory.

```
mkdir compare_repo
cd compare_repo
git init

cp -r ../xv6-original/* .
git add .
git commit -m "original"

cp -r ../xv6-modified/* .
git add .
git commit -m "modified"

git diff HEAD~1 HEAD '*.c' '*.h' > ../diff_report.txt
```